

Mel-Frequency Ceptral Coeffienents(MFCC) feature extraction for Sound Classification

For sound classification like the [Cornell Birdcall Identification \(https://www.kaggle.com/c/birdsong-recognition/overview\)](https://www.kaggle.com/c/birdsong-recognition/overview) is usually using the MFCC feature.

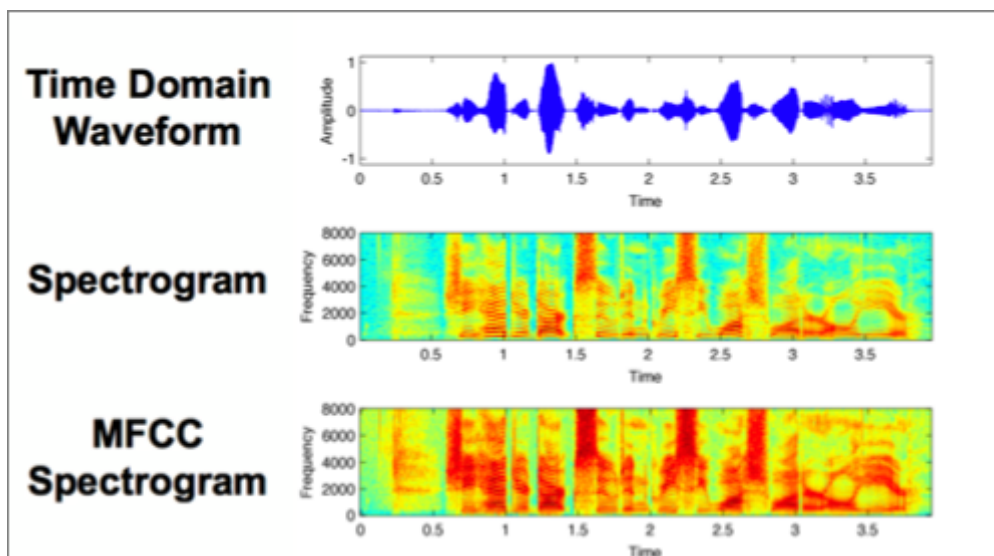
It takes few hours for Cornell Birdcall Identification datasets. I will share extracted feature as dataset after the execution in colab.

In this notebook, I just use 3 mp3 files for each bird class. (check the LIMIT variable)

Please enjoy it and don't forget to vote it.

Feel free to give an advice.

Mel-Frequency Cepstral Coefficients (MFCCs)

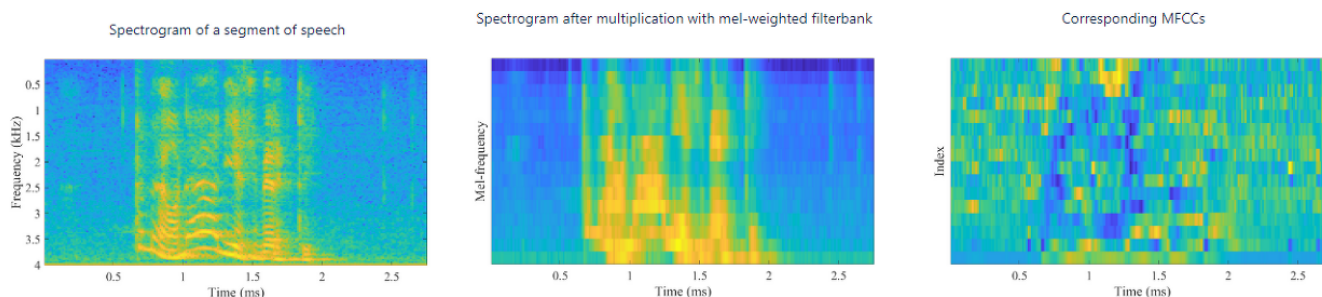


The log-spectrum already takes into account perceptual sensitivity on the magnitude axis, by expressing magnitudes on the logarithmic-axis. The other dimension is then the frequency axis.

There exists a multitude of different criteria with which to quantify accuracy on the frequency scale and there are, correspondingly, a multitude of perceptually motivated frequency scales including the equivalent rectangular bandwidth (ERB) scale, the Bark scale, and the mel-scale. Probably through an arbitrary choice mainly due to tradition, in this context we will focus on the mel-scale. This scale describes the perceptual distance between pitches of different frequencies.

Though the argumentation for the MFCCs is not without problems, it has become the most used feature in speech and audio recognition applications. It is used because it works and because it has relatively low complexity and it is straightforward to implement. Simply stated,

if you're unsure which inputs to give to a speech and audio recognition engine, try first the MFCCs.



The beneficial properties of the MFCCs include:

Quantifies the gross-shape of the spectrum (the spectral envelope), which is important in, for example, identification of vowels. At the same time, it removes fine spectral structure (micro-level structure), which is often less important. It thus focuses on that part of the signal which is typically most informative. Straightforward and computationally reasonably efficient calculation. Their performance is well-tested and -understood. Some of the issues with the MFCC include:

The choice of perceptual scale is not well-motivated. Scales such as the ERB or gamma-tone filterbanks might be better suited. However, these alternative filterbanks have not demonstrated consistent benefit, whereby the mel-scale has persisted. MFCCs are not robust to noise. That is, the performance of MFCCs in presence of additive noise, in comparison to other features, has not always been good. The choice of triangular weighting filters $w_{k,h}$ is arbitrary and not based on well-grounded motivations. Alternatives have been presented, but they have not gained popularity, probably due to minor effect on outcome. The MFCCs work well in analysis but for synthesis, they are problematic. Namely, it is difficult to find an inverse transform (from MFCCs to power spectra) which is simultaneously unbiased (=accurate) and congruent with its physical representation (=power spectrum must be positive).

ref: <https://wiki.aalto.fi/display/ITSP/Cepstrum+and+MFCC>
<https://wiki.aalto.fi/display/ITSP/Cepstrum+and+MFCC>

In [1]:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import os
import glob
import librosa
import librosa.display
from tqdm import tqdm_notebook as tqdm
from keras.models import Model
from keras.utils import np_utils

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
import matplotlib.pyplot as plt
```

Using TensorFlow backend.

In [2]:

```
LIMIT = 3
```

In [3]:

```
!ls ../input/birdsong-recognition
```

```
example_test_audio          sample_submission.csv  train_audio
example_test_audio_metadata.csv test.csv
example_test_audio_summary.csv train.csv
```

In [4]:

```
df_train = pd.read_csv('../input/birdsong-recognition/train.csv')  
df_train
```

Out[4]:

	rating	playback_used	ebird_code	channels	date	pitch	duration	filename
0	3.5	no	aldfly	1 (mono)	2013-05-25	Not specified	25	XC1348
1	4.0	no	aldfly	2 (stereo)	2013-05-27	both	36	XC1354
2	4.0	no	aldfly	2 (stereo)	2013-05-27	both	39	XC1354
3	3.5	no	aldfly	2 (stereo)	2013-05-27	both	33	XC1354
4	4.0	no	aldfly	2 (stereo)	2013-05-27	both	36	XC1354
...
21370	4.5	no	yetvir	1 (mono)	2019-05-15	both	28	XC4776
21371	3.5	no	yetvir	1 (mono)	2017-05-14	Not specified	52	XC5003
21372	5.0	no	yetvir	1 (mono)	2017-06-10	Not specified	96	XC5012
21373	3.5	no	yetvir	2 (stereo)	2009-05-06	level	35	XC5482
21374	3.5	no	yetvir	2 (stereo)	2010-06-09	level	103	XC5576

21375 rows × 35 columns

In [5]:

```
!ls ../input/birdsong-recognition/train_audio

train_dir = '../input/birdsong-recognition/train_audio'
test_idr = '../input/birdsong-recognition/test_audio'
```

aldfly	bktspa	canwre	easkin	hergul	magwar	pinsis	sa
vspa	wesmea						
ameavo	blkpho	carwre	easmea	herthr	mallar3	pinwar	sa
ypho	wessan						
amebit	blugrb1	casfin	easpho	hoomer	marwre	plsvir	sc
atan	westan						
amecro	blujay	caster1	eastow	hoowar	merlin	prawar	sc
oori	wewpew						
amegfi	bnhcow	casvir	eawpew	horgre	moublu	purfin	se
mplo	whbnut						
amekes	boboli	cedwax	eucdov	horlar	mouchi	pygnut	se
msan	whcspa						
amepip	bongul	chispa	eursta	houfin	moudov	rebmer	sh
eowl	whfibi						
amered	brdowl	chiswi	evegro	houspa	norcar	rebnut	sh
shaw	whtspa						
amerob	brebla	chswar	fiespa	houwre	norfli	rebsap	sn
obun	whtswi						
amewig	brespa	chukar	fiscro	indbun	norhar2	rebwoo	sn
ogoo	wilfly						
amewoo	brncre	clanut	foxspa	juntit1	normoc	redcro	so
lsan	wilsni1						
amtspa	brnthr	cliswa	gadwal	killde	norpar	redhea	so
nspa	wiltur						
annhum	brthum	comgol	gcrfin	labwoo	norpin	reevir1	so
ra	winwre3						
astfly	brwhaw	comgra	gnttow	larspa	norsho	renpha	sp
osan	wlswar						
baisan	btbwar	comloo	gnwtea	lazbun	norwat	reshaw	sp
otow	wooduc						
baleag	btnwar	commer	gockin	leabit	nrswa	rethaw	st
ejay	wooscj2						
balori	btywar	comnig	gocspa	leafly	nutwoo	rewbla	sw
ahaw	woothr						
banswa	buffle	comrav	goleag	leasan	olsfly	ribgul	sw
aspa	y00475						
barswa	buggna	comred	grbher3	lecthr	orcwar	rinduc	sw
athr	yebfly						
bawwar	buhvir	comter	grcfly	lesgol	osprey	robgro	tr
eswa	yebsap						
belkin1	bulori	comyel	greegr	lesnig	ovenbi1	rocpig	tr
uswa	yehbla						
belspa2	bushti	coohaw	greroa	lesyel	palwar	rocwre	tu
ftit	yelwar						

bewwre	buwtea	coshum	greyel	lewwoo	pasfly	rthhum	tu
nswa	yerwar						
bkbuc	buwwar	cowscj1	grhowl	linspa	pecsan	ruckin	ve
ery	yetvir						
bkbmag1	cacwre	daejun	grnher	lobcur	perfal	rudduc	ve
sspa							
bkbwar	calgul	doccor	grtgra	lobdow	phaino	rufgro	vi
gswa							
bkcchi	calqua	dowwoo	grycat	logshr	pibgre	rufhum	wa
rvir							
bkchum	camwar	dusfly	gryfly	lotduc	pilwoo	rusbla	we
sblu							
bkhgro	cangoo	eargre	haiwoo	louwat	pingro	sagspa1	we
sgre							
bkpwar	canwar	easblu	hamfly	macwar	pinjay	sagthr	we
skin							

Extract Feature using MFCC()

In [6]:

```
def mfcc_extract(filename):
    try:
        y, sr = librosa.load(filename, sr = 44100)
        mfcc = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=13, n_fft=int(0.02*sr), hop_length=int(0.01*sr))
        return mfcc
    except:
        return
```

In [7]:

```
def parse_audio_files(parent_dir, sub_dirs, limit):
    labels = []
    features = []
    for label, sub_dir in enumerate(tqdm(sub_dirs)):
        i = 0
        for fn in glob.glob(os.path.join(parent_dir, sub_dir, "*.mp3")):
            if i >= limit:
                break
            features.append(mfcc_extract(fn))
            labels.append(label)
            i+=1
    return features, labels
```

In [8]:

```
%%time

train_cat_dirs = glob.glob(train_dir+'/*')
train_cat = []
for cat_dir in train_cat_dirs:
    tmp = cat_dir.split('/')[1]
    train_cat.append(tmp)
print('the number of kinds:', len(train_cat))

class_num = len(train_cat)
features, labels = parse_audio_files(train_dir, train_cat, LIMIT)
```

the number of kinds: 264

100%

264/264 [19:57<00:00, 4.53s/it]

CPU times: user 19min 34s, sys: 2min 3s, total: 21min 38s

Wall time: 19min 57s

In [9]:

```
print(len(features))
print(features[0].shape)
```

```
792
(13, 8164)
```

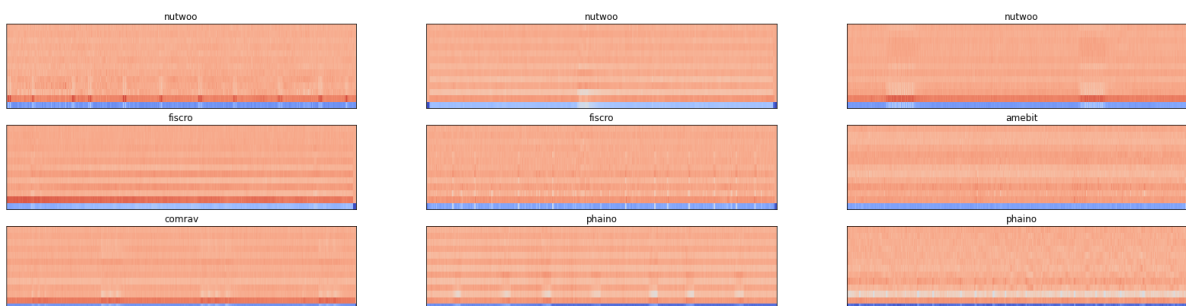
In [10]:

```
# plot few features

fig = plt.figure(figsize=(28,24))
for i,mfcc in enumerate(tqdm(features[:100])):
    if i%40 < 3 :
        sub = plt.subplot(10,3,i%40+3*(i/40)+1)
        librosa.display.specshow(mfcc,vmin=-700,vmax=300)
        if ((i%40+3*(i/40)+1)%3==0) :
            plt.colorbar()
        sub.set_title(train_cat[labels[i]])
plt.show()
```

100%

100/100 [00:00<00:00, 302.42it/s]



In [11]:

```
df_submission = pd.read_csv('../input/birdsong-recognition/sample_submission.csv')
df_submission.to_csv('submission.csv', index = None)
```